

Classification des Caractères Tifinagh (niveaux de gris) avec un Réseau de Neurones Multiclasses

Pr. M. BENADDY
Masters : IMSD & IAA

2024-2025

Ce TP consiste à implémenter un réseau de neurones multiclasses pour la classification des caractères Tifinagh générés à partir de la base de données AMHCD. La base de données contient 28182 images (64x64 pixels) réparties sur les 33 classes qui représente l'alphabet Tifinagh. Le modèle à implémenter en Python, qui est une adaptation du TP de la classification binaire, utilise un perceptron multicouche (MLP) avec deux couches cachées (paramétrable), des activations ReLU, et une sortie softmax. Vous devez décrire les formules mathématiques sous-jacentes.

1 Introduction

La reconnaissance des caractères manuscrits est un défi important en vision par ordinateur, particulièrement pour les scripts moins étudiés comme le Tifinagh, utilisé par les communautés amazighes. Dans ce projet vous devez implémenter un réseau de neurones multiclasses pour classifier ces caractères, en utilisant un MLP avec des couches de 64 et 32 neurones cachés, entraîné sur des images redimensionnées à 32x32 pixels et aplaties en vecteurs de 1024 caractéristiques.

Télécharger la base de données : https://drive.google.com/file/d/1g03sIzi8F855KRMi0wmJesc/view?usp=share_link

2 Formules Mathématiques

Les formules clés du modèle.

2.1 L'opération forward

Pour une couche l , la sortie linéaire et l'activation sont :

$$Z^{[l]} = A^{[l-1]}W^{[l]} + b^{[l]} \quad (1)$$

$$A^{[l]} = g^{[l]}(Z^{[l]}) \quad (2)$$

où $A^{[0]} = X$, $X \in \mathbb{R}^{m \times 1024}$ est l'entrée, $W^{[l]} \in \mathbb{R}^{n^{[l-1]} \times n^{[l]}}$ les poids, et $b^{[l]} \in \mathbb{R}^{1 \times n^{[l]}}$ les biais. Pour les couches cachées ($l = 1, 2$) :

$$g^{[l]}(z) = \text{ReLU}(z) = \max(0, z) \quad (3)$$

Pour la couche de sortie ($l = 3$) :

$$A^{[3]} = \text{softmax}(Z^{[3]}), \quad \hat{y}_{i,c} = \frac{e^{z_{i,c}}}{\sum_{j=1}^{33} e^{z_{i,j}}} \quad (4)$$

2.2 La perte

L'entropie :

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{c=1}^{33} y_{i,c} \log(\hat{y}_{i,c}) \quad (5)$$

où $y_{i,c}$ est 1 si la classe c est correcte, 0 sinon.

2.3 Précision

La précision est :

$$\text{Accuracy} = \frac{1}{m} \sum_{i=1}^m (\arg \max_c (\hat{y}_{i,c}) = \arg \max_c (y_{i,c})) \quad (6)$$

2.4 Rétropropagation

Le gradient initial pour la couche de sortie est :

$$\frac{\partial J}{\partial Z^{[3]}} = \hat{y} - Y \quad (7)$$

Pour les couches cachées ($l = 2, 1$) :

$$\frac{\partial J}{\partial Z^{[l]}} = \left(\frac{\partial J}{\partial Z^{[l+1]}} W^{[l+1]T} \right) \cdot \text{ReLU}'(Z^{[l]}) \quad (8)$$

où :

$$\text{ReLU}'(z) = \begin{cases} 1 & \text{si } z > 0 \\ 0 & \text{sinon} \end{cases} \quad (9)$$

Les gradients des paramètres sont :

$$\frac{\partial J}{\partial W^{[l]}} = \frac{1}{m} (A^{[l-1]})^T \cdot \frac{\partial J}{\partial Z^{[l]}} \quad (10)$$

$$\frac{\partial J}{\partial b^{[l]}} = \frac{1}{m} \sum_{i=1}^m \frac{\partial J}{\partial Z^{[l]}} \quad (11)$$

Mise à jour des paramètres :

$$W^{[l]} \leftarrow W^{[l]} - \eta \frac{\partial J}{\partial W^{[l]}} \quad (12)$$

$$b^{[l]} \leftarrow b^{[l]} - \eta \frac{\partial J}{\partial b^{[l]}} \quad (13)$$

où $\eta = 0.01$.

3 Implémentation

Le code suivant implémente le modèle en Python, avec le chargement des données de Tifinagh partagée, le prétraitement, l'entraînement, et l'évaluation.

```
1 import os
2 import pandas as pd
3 import numpy as np
4 import cv2
5 from sklearn.model_selection import train_test_split
6 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
7 from sklearn.metrics import confusion_matrix,
8     classification_report
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11
12 # Fonctions d'activation
13 def relu(x):
14     """
15     ReLU activation: max(0, x)
16     """
17     assert isinstance(x, np.ndarray), "Input to ReLU must be a
18         numpy array"
19     # TODO
20     assert np.all(result >= 0), "ReLU output must be non-negative"
21     return result
22
23 def relu_derivative(x):
24     """
25     Derivative of ReLU: 1 if x > 0, else 0
26     """
27     assert isinstance(x, np.ndarray), "Input to ReLU derivative
28         must be a numpy array"
29     result = # TODO
30     assert np.all((result == 0) | (result == 1)), "ReLU derivative
31         must be 0 or 1"
32     return result
33
34 def softmax(x):
35     """
36     Softmax activation: exp(x) / sum(exp(x))
37     """
38     assert isinstance(x, np.ndarray), "Input to softmax must be a
39         numpy array"
40     exp_x = # TODO
41     result = exp_x / np.sum(exp_x, axis=1, keepdims=True)
42     assert np.all((result >= 0) & (result <= 1)), "Softmax output
43         must be in [0, 1]"
44     assert np.allclose(np.sum(result, axis=1), 1), "Softmax output
45         must sum to 1 per sample"
46     return result
```

```

41 # Classe MultiClassNeuralNetwork
42 class MultiClassNeuralNetwork:
43     def __init__(self, layer_sizes, learning_rate=0.01):
44         """
45         Initialize the neural network with given layer sizes and
46         learning rate.
47         layer_sizes: List of integers [input_size, hidden1_size,
48         ..., output_size]
49         """
50         assert isinstance(layer_sizes, list) and len(layer_sizes)
51             >= 2, "layer_sizes must be a list with at least 2
52             elements"
53         assert all(isinstance(size, int) and size > 0 for size in
54             layer_sizes), "All layer sizes must be positive
55             integers"
56         assert isinstance(learning_rate, (int, float)) and
57             learning_rate > 0, "Learning rate must be a positive
58             number"
59
60         self.layer_sizes = layer_sizes
61         self.learning_rate = learning_rate
62         self.weights = []
63         self.biases = []
64
65         # Initialisation des poids et biais
66         np.random.seed(42)
67         for i in range(len(layer_sizes) - 1):
68             w = np.random.randn(layer_sizes[i], layer_sizes[i+1])
69                 * 0.01
70             b = np.zeros((1, layer_sizes[i+1]))
71             assert w.shape == (layer_sizes[i], layer_sizes[i+1]),
72                 f"Weight matrix {i+1} has incorrect shape"
73             assert b.shape == (1, layer_sizes[i+1]), f"Bias vector
74                 {i+1} has incorrect shape"
75             self.weights.append(w)
76             self.biases.append(b)
77
78     def forward(self, X):
79         """
80         Forward propagation:  $Z^{\{[1]\}} = A^{\{[1-1]\}} W^{\{[1]\}} + b^{\{[1]\}}$ 
81          $A^{\{[1]\}} = g(Z^{\{[1]\}})$ 
82         """
83         assert isinstance(X, np.ndarray), "Input X must be a numpy
84             array"
85         assert X.shape[1] == self.layer_sizes[0], f"Input
86             dimension ({X.shape[1]}) must match input layer size ({
87             self.layer_sizes[0]})"
88
89         self.activations = [X]
90         self.z_values = []

```

```

77     for i in range(len(self.weights) - 1):
78         z = # TODO
79         assert z.shape == (X.shape[0], self.layer_sizes[i+1]),
            f"Z^{[i+1]} has incorrect shape"
80         self.z_values.append(z)
81         self.activations.append(relu(z))
82
83     z = # TODO
84     assert z.shape == (X.shape[0], self.layer_sizes[-1]), "
        Output Z has incorrect shape"
85     self.z_values.append(z)
86     output = # TODO softmax call
87     assert output.shape == (X.shape[0], self.layer_sizes[-1]),
        "Output A has incorrect shape"
88     self.activations.append(output)
89
90     return self.activations[-1]
91
92 def compute_loss(self, y_true, y_pred):
93     """
94     Categorical Cross-Entropy:  $J = -1/m * \sum(y\_true * \log($ 
95          $y\_pred))$ 
96     """
97     assert isinstance(y_true, np.ndarray) and isinstance(
98         y_pred, np.ndarray), "Inputs to loss must be numpy
99         arrays"
100     assert y_true.shape == y_pred.shape, "y_true and y_pred
101         must have the same shape"
102
103     y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
104     loss = # TODO
105     assert not np.isnan(loss), "Loss computation resulted in
106         NaN"
107     return loss
108
109 def compute_accuracy(self, y_true, y_pred):
110     """
111     Compute accuracy: proportion of correct predictions
112     """
113     assert isinstance(y_true, np.ndarray) and isinstance(
114         y_pred, np.ndarray), "Inputs to accuracy must be numpy
115         arrays"
116     assert y_true.shape == y_pred.shape, "y_true and y_pred
117         must have the same shape"
118
119     predictions = # TODO
120     true_labels = # TODO
121     accuracy = # TODO
122     assert 0 <= accuracy <= 1, "Accuracy must be between 0 and
123         1"
124     return accuracy

```

```

116
117 def backward(self, X, y, outputs):
118     """
119     Backpropagation: compute  $dW^{[l]}$ ,  $db^{[l]}$  for each layer
120     """
121     assert isinstance(X, np.ndarray) and isinstance(y, np.
122         ndarray) and isinstance(outputs, np.ndarray), "Inputs
123         to backward must be numpy arrays"
124     assert X.shape[1] == self.layer_sizes[0], f"Input
125         dimension ({X.shape[1]}) must match input layer size ({
126         self.layer_sizes[0]})"
127     assert y.shape == outputs.shape, "y and outputs must have
128         the same shape"
129
130     m = X.shape[0]
131     self.d_weights = # TODO
132     self.d_biases = # TODO
133
134     dZ = outputs - y # Gradient pour softmax + cross-entropy
135     assert dZ.shape == outputs.shape, "dZ for output layer has
136         incorrect shape"
137     self.d_weights[-1] = (self.activations[-2].T @ dZ) / m
138     self.d_biases[-1] = np.sum(dZ, axis=0, keepdims=True) / m
139
140     for i in range(len(self.weights) - 2, -1, -1):
141         dZ = #TODO
142         assert dZ.shape == (X.shape[0], self.layer_sizes[i+1])
143             , f"dZ^{[i+1]} has incorrect shape"
144         self.d_weights[i] = #TODO
145         self.d_biases[i] = #TODO
146
147     # TODO: Ajouter une regularisation L2 aux gradients des
148         poids
149     #  $dW^{[l]} += \lambda * W^{[l]} / m$ , o  $\lambda$  est le
150         coefficient de regularisation
151
152     for i in range(len(self.weights)):
153         self.weights[i] -= self.learning_rate * self.d_weights
154             [i]
155         self.biases[i] -= self.learning_rate * self.d_biases[i]
156
157 def train(self, X, y, X_val, y_val, epochs, batch_size):
158     """
159     Train the neural network using mini-batch SGD, with
160     validation
161     """
162     assert isinstance(X, np.ndarray) and isinstance(y, np.
163         ndarray), "X and y must be numpy arrays"
164     assert isinstance(X_val, np.ndarray) and isinstance(y_val,
165         np.ndarray), "X_val and y_val must be numpy arrays"

```

```

153     assert X.shape[1] == self.layer_sizes[0], f"Input
        dimension ({X.shape[1]}) must match input layer size ({
        self.layer_sizes[0]})"
154     assert y.shape[1] == self.layer_sizes[-1], f"Output
        dimension ({y.shape[1]}) must match output layer size
        ({self.layer_sizes[-1]})"
155     assert X_val.shape[1] == self.layer_sizes[0], f"Validation
        input dimension ({X_val.shape[1]}) must match input
        layer size ({self.layer_sizes[0]})"
156     assert y_val.shape[1] == self.layer_sizes[-1], f"
        Validation output dimension ({y_val.shape[1]}) must
        match output layer size ({self.layer_sizes[-1]})"
157     assert isinstance(epochs, int) and epochs > 0, "Epochs
        must be a positive integer"
158     assert isinstance(batch_size, int) and batch_size > 0, "
        Batch size must be a positive integer"
159
160     train_losses = []
161     val_losses = []
162     train_accuracies = []
163     val_accuracies = []
164
165     for epoch in range(epochs):
166         indices = np.random.permutation(X.shape[0])
167         X_shuffled = X[indices]
168         y_shuffled = y[indices]
169
170         epoch_loss = 0
171         for i in range(0, X.shape[0], batch_size):
172             X_batch = X_shuffled[i:i+batch_size]
173             y_batch = y_shuffled[i:i+batch_size]
174
175             outputs = self.forward(X_batch)
176             epoch_loss += self.compute_loss(y_batch, outputs)
177             self.backward(X_batch, y_batch, outputs)
178
179         # Calculer les pertes et accuracies
180         train_loss = epoch_loss / (X.shape[0] // batch_size)
181         train_pred = self.forward(X)
182         train_accuracy = self.compute_accuracy(y, train_pred)
183         val_pred = self.forward(X_val)
184         val_loss = self.compute_loss(y_val, val_pred)
185         val_accuracy = self.compute_accuracy(y_val, val_pred)
186
187         train_losses.append(train_loss)
188         val_losses.append(val_loss)
189         train_accuracies.append(train_accuracy)
190         val_accuracies.append(val_accuracy)
191
192         if epoch % 10 == 0:
193             print(f"Epoch {epoch}, Train Loss: {train_loss:.4f}

```

```

194         }, Val Loss: {val_loss:.4f}, "
195         f"Train Acc: {train_accuracy:.4f}, Val Acc:
196         {val_accuracy:.4f}")
197
198     return train_losses, val_losses, train_accuracies,
199         val_accuracies
200
201     def predict(self, X):
202         """
203         Predict class labels
204         """
205         assert isinstance(X, np.ndarray), "Input X must be a numpy
206             array"
207         assert X.shape[1] == self.layer_sizes[0], f"Input
208             dimension ({X.shape[1]}) must match input layer size ({
209             self.layer_sizes[0]})"
210
211         outputs = self.forward(X)
212         predictions = np.argmax(outputs, axis=1)
213         assert predictions.shape == (X.shape[0],), "Predictions
214             have incorrect shape"
215         return predictions
216
217     # D finir le chemin vers le dossier d compress
218     data_dir = os.path.join(os.getcwd(), 'amhcd-data-64/tifinagh-
219         images/')
220     print(data_dir)
221     current_working_directory = os.getcwd()
222     print(current_working_directory)
223
224     # Charger le fichier CSV contenant les tiquettes
225     try:
226         labels_df = pd.read_csv(os.path.join(data_dir, 'amhcd-data-64/
227             labels-map.csv'))
228         assert 'image_path' in labels_df.columns and 'label' in
229             labels_df.columns, "CSV must contain 'image_path' and '
230             label' columns"
231     except FileNotFoundError:
232         print("labels-map.csv not found. Please check the dataset
233             structure.")
234
235     # Alternative : construire un DataFrame partir des dossiers
236     image_paths = []
237     labels = []
238     for label_dir in os.listdir(data_dir):
239         label_path = os.path.join(data_dir, label_dir)
240         if os.path.isdir(label_path):
241             for img_name in os.listdir(label_path):
242                 image_paths.append(os.path.join(label_path,
243                     img_name))
244                 labels.append(label_dir)
245     labels_df = pd.DataFrame({'image_path': image_paths, 'label':

```



```

        labels}))
232
233 # Vérifier le DataFrame
234 assert not labels_df.empty, "No data loaded. Check dataset files."
235 print(f"Loaded {len(labels_df)} samples with {labels_df['label']}.
        nunique()} unique classes.")
236
237 # Encoder les étiquettes
238 label_encoder = LabelEncoder()
239 labels_df['label_encoded'] = label_encoder.fit_transform(labels_df
        ['label'])
240 num_classes = len(label_encoder.classes_)
241
242 # Fonction pour charger et prétraiter une image
243 def load_and_preprocess_image(image_path, target_size=(32, 32)):
244     """
245     Load and preprocess an image: convert to grayscale, resize,
        normalize
246     """
247     assert os.path.exists(image_path), f"Image not found: {
        image_path}"
248     img = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
249     assert img is not None, f"Failed to load image: {image_path}"
250     img = cv2.resize(img, target_size)
251     img = img.astype(np.float32) / 255.0 # Normalisation
252     return img.flatten() # Aplatir pour le réseau de neurones
253
254 # Charger toutes les images
255 X = np.array([load_and_preprocess_image(os.path.join(data_dir,
        path)) for path in labels_df['image_path']])
256 y = labels_df['label_encoded'].values
257
258 # Vérifier les dimensions
259 assert X.shape[0] == y.shape[0], "Mismatch between number of
        images and labels"
260 assert X.shape[1] == 32 * 32, f"Expected flattened image size of
        {32*32}, got {X.shape[1]}"
261
262 # Diviser en ensembles d'entraînement, validation et test
263 X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size
        =0.2, stratify=y, random_state=42)
264 X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp,
        test_size=0.25, stratify=y_temp, random_state=42)
265
266 # Convertir explicitement en NumPy arrays
267 X_train = np.array(X_train)
268 X_val = np.array(X_val)
269 X_test = np.array(X_test)
270 y_train = np.array(y_train)
271 y_val = np.array(y_val)
272 y_test = np.array(y_test)

```

```

273
274 assert X_train.shape[0] + X_val.shape[0] + X_test.shape[0] == X.
    shape[0], "Train-val-test split sizes must sum to total samples
    "
275
276 print(f"Train: {X_train.shape[0]} samples, Validation: {X_val.
    shape[0]} samples, Test: {X_test.shape[0]} samples")
277
278 # Encoder les tiquettes en one-hot pour la classification
    multiclasse
279 one_hot_encoder = OneHotEncoder(sparse_output=False)
280 y_train_one_hot = np.array(one_hot_encoder.fit_transform(y_train.
    reshape(-1, 1)))
281 y_val_one_hot = np.array(one_hot_encoder.transform(y_val.reshape
    (-1, 1)))
282 y_test_one_hot = np.array(one_hot_encoder.transform(y_test.reshape
    (-1, 1)))
283
284 # V rifier que les tableaux one-hot sont des NumPy arrays
285 assert isinstance(y_train_one_hot, np.ndarray), "y_train_one_hot
    must be a numpy array"
286 assert isinstance(y_val_one_hot, np.ndarray), "y_val_one_hot must
    be a numpy array"
287 assert isinstance(y_test_one_hot, np.ndarray), "y_test_one_hot
    must be a numpy array"
288
289 # Cr er et entra ner le mod le
290 layer_sizes = [X_train.shape[1], 64, 32, num_classes] # 64 et 32
    neurones cach s, 33 classes
291 nn = MultiClassNeuralNetwork(layer_sizes, learning_rate=0.01)
292 train_loss, val_loss, train_acc, val_acc = nn.
    train(
293     X_train, y_train_one_hot, X_val, y_val_one_hot, epochs=100,
        batch_size=32
294 )
295
296 # TODO: Ajouter une validation crois e pour valuer la
    robustesse du mod le
297 # TODO: Impl menter l'optimiseur Adam pour une meilleure
    convergence
298
299 # Pr dictions et valuation
300 y_pred = nn.predict(X_test)
301 print("\nRapport de classification (Test set) :")
302 print(classification_report(y_test, y_pred, target_names=
    label_encoder.classes_))
303
304 # Matrice de confusion
305 cm = confusion_matrix(y_test, y_pred)
306 plt.figure(figsize=(10, 8))
307 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

```

```

308 plt.title('Matrice de confusion (Test set)')
309 plt.xlabel('Pr dit')
310 plt.ylabel('R el')
311 plt.savefig('confusion_matrix.png')
312 plt.close()
313
314 # Courbes de perte et d'accuracy
315 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
316
317 # Courbe de perte
318 ax1.plot(train_lossess, label='Train Loss')
319 ax1.plot(val_lossess, label='Validation Loss')
320 ax1.set_title('Courbe de perte')
321 ax1.set_xlabel(' poque ')
322 ax1.set_ylabel('Perte')
323 ax1.legend()
324
325 # Courbe d'accuracy
326 ax2.plot(train_accuracies, label='Train Accuracy')
327 ax2.plot(val_accuracies, label='Validation Accuracy')
328 ax2.set_title('Courbe de pr cision')
329 ax2.set_xlabel(' poque ')
330 ax2.set_ylabel('Pr cision')
331 ax2.legend()
332
333 plt.tight_layout()
334 fig.savefig('loss_accuracy_plot.png')
335 plt.close()

```

4 Visualisations

Les résultats atteints pour la configuration suivante : $32 \times 32, 64, 32, 33$

```

1 layer_sizes = [X_train.shape[1], 64, 32, num_classes] # 64 et 32
   neurones cach s, 33 classes
2 nn = MultiClassNeuralNetwork(layer_sizes, learning_rate=0.01)
3 train_lossess, val_lossess, train_accuracies, val_accuracies = nn.
   train(
4     X_train, y_train_one_hot, X_val, y_val_one_hot, epochs=100,
       batch_size=32
5 )

```

5 Travail à faire

En plus de l'implémentation du modèle vous devez rédiger un rapport (article en Anglais) selon le format IMRAD déjà évoquée dans le TP de classification binaire.

NB : Le lien vers le code github doit être inclu dans le rapport.

Des améliorations incluent (Bonus) :




Figure 1: Matrice de confusion pour l'ensemble de test, montrant les prédictions correctes (diagonale) et les erreurs.

- **Régularisation L2** : Ajouter $\frac{\lambda}{m}W^{[l]}$ aux gradients.
- **Optimiseur Adam** : Utiliser des moments adaptatifs pour une meilleure convergence.
- **Validation croisée** : Évaluer la robustesse avec K-fold.
- **Augmentation des données** : Appliquer des rotations et translations.
- **Reprendre le travail avec la base de données AMHCD citée dans la bibliographie**

References

- [1] Benaddy, M., et al. "Amazigh Handwritten Character Database (AMHCD)." Kaggle, 2020. Disponible sur : <https://www.kaggle.com/datasets/benaddym/amazigh-handwritten-character-database-amhcd>.




Figure 2: Courbes de perte (gauche) et de précision (droite) pour les ensembles d'entraînement et de validation.